

Assessing Performance Evolution for Configurable Systems

A Methodology



Advisors: Prof. Dr.-Ing. Ina Schaefer, Prof. Dr.-Ing. Norbert Siegmund

January 23, 2018
Stefan Mühlbauer

TECH —

Here's how, and why, the Spectre and Meltdown patches will hurt performance

Now that microcode and patches are starting to ship, a clearer picture is emerging.

PETER BRIGHT - 1/11/2018, 10:30 PM

Apple Confirms It Degrades Your Old iPhone's Performance



Ewan Spence, CONTRIBUTOR
[FULL BIO](#) ✓

Opinions expressed by Forbes Contributors are their own.

Intel-Benchmarks zu Meltdown/Spectre: Performance sackt um bis zu 10 Prozent ab, SSD-I/O deutlich mehr



11.01.2018 12:15 Uhr – Martin Fischer

Apple: Decreased performance with old batteries is a feature, not a bug




Chris Mills [@chrisfmills](#)
December 28th, 2017 at 6:04 PM

Share


Tweet

Configurable Software Systems



Configurable Software Systems

- Configurability: Compile-time vs load-time



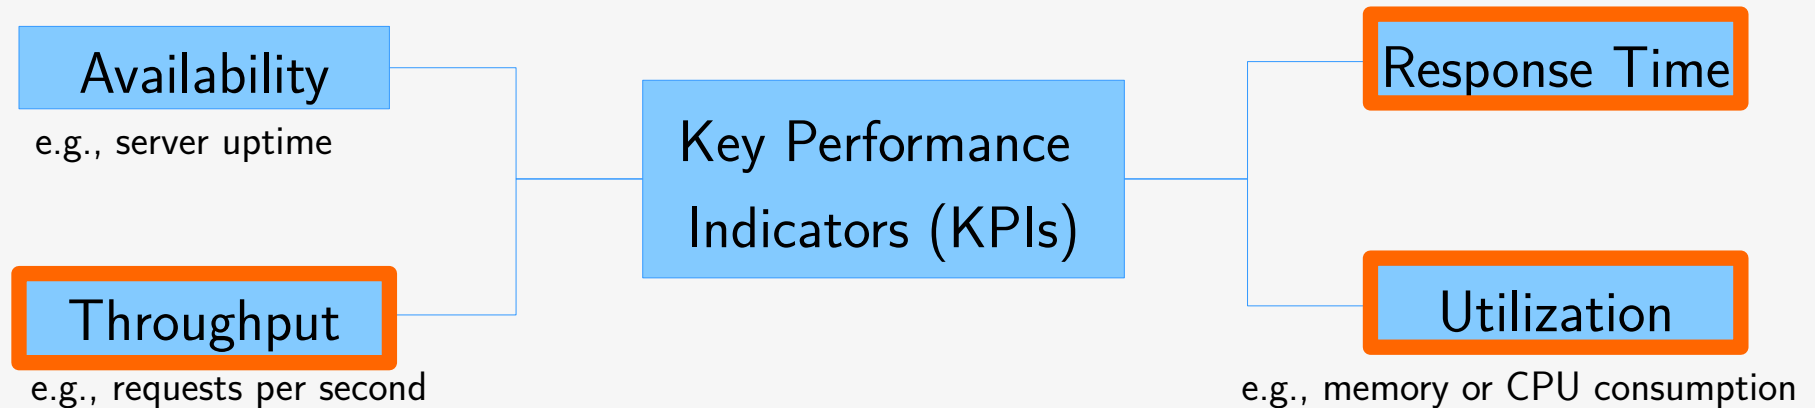
```
~$ zip -e -9 file.txt
```

- Unanticipated behavior can emerge with selections of multiple features (feature interaction)
 - Example: Compression and Encryption

*Encrypting compressed data can be **faster** than encrypting raw data, since compressed data is already more compact.*


Software Performance Metrics

- Software performance: How **efficiently** is a task executed?
- Performance footprints can be outlined by key performance indicators:



Software Evolution

- Evolution: adaption to changing contexts and requirements
- Software can *grow more complex* over time as it evolves:



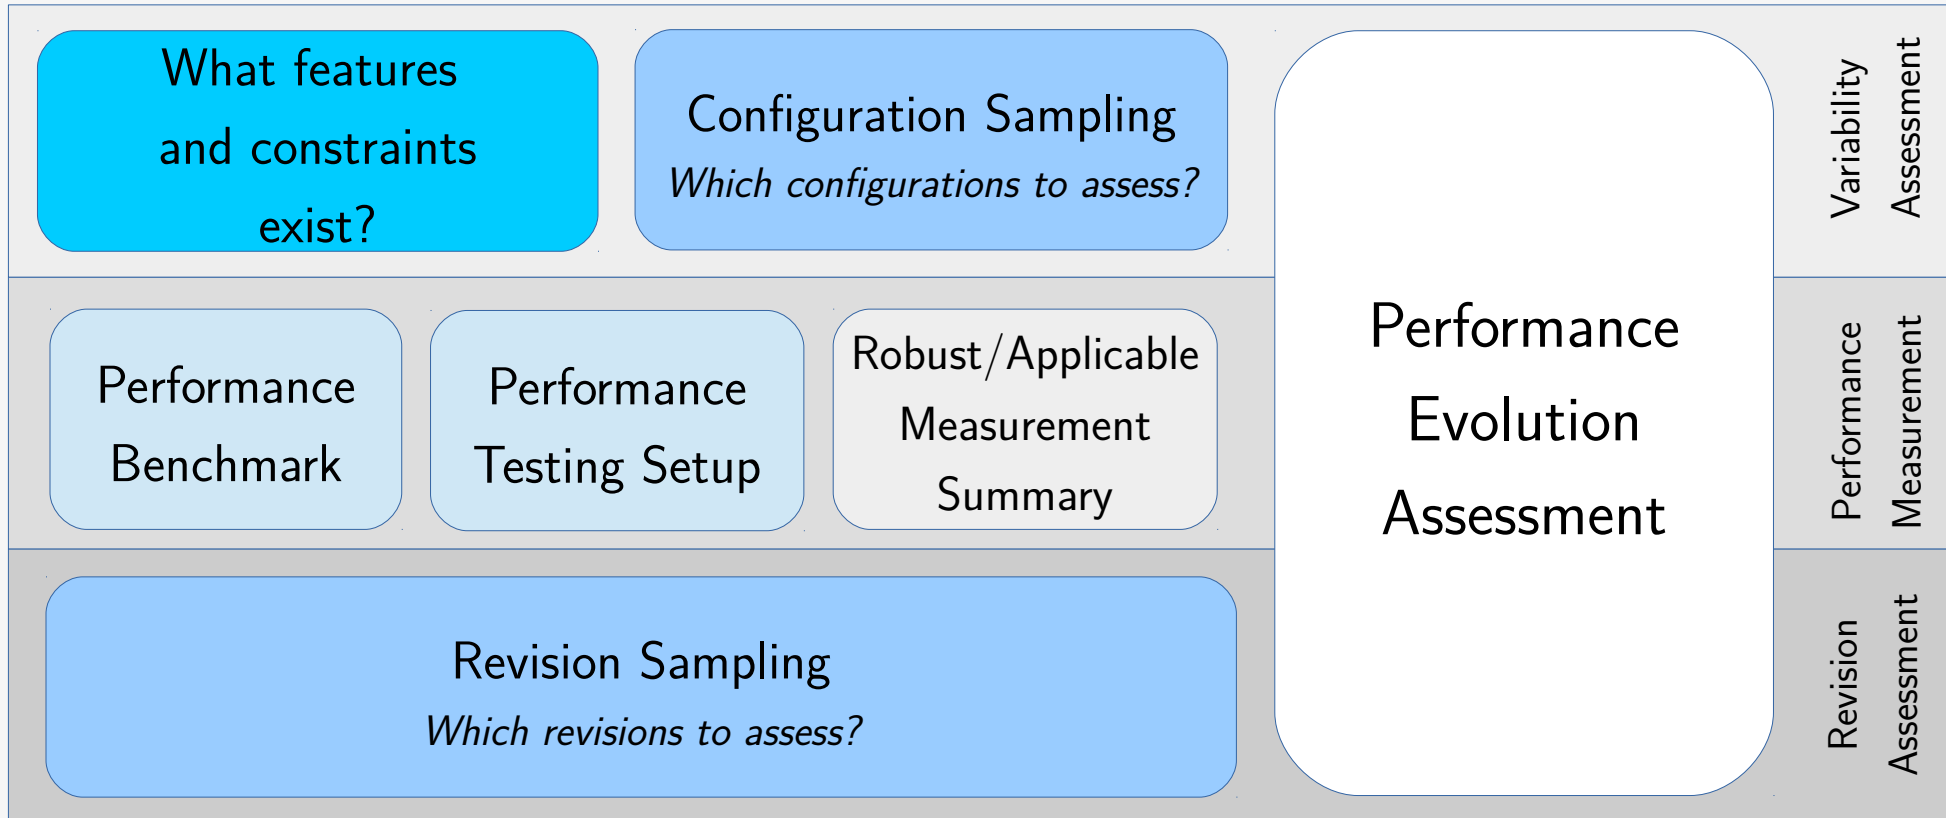
Adopted: Fig. 1 from Israeli, A., & Feitelson, D. G. (2010). *The Linux kernel as a case study in software evolution*. Journal of Systems and Software, 83(3), 485-501.

- Software can “erode” as it evolves, leading to degradation of overall software quality

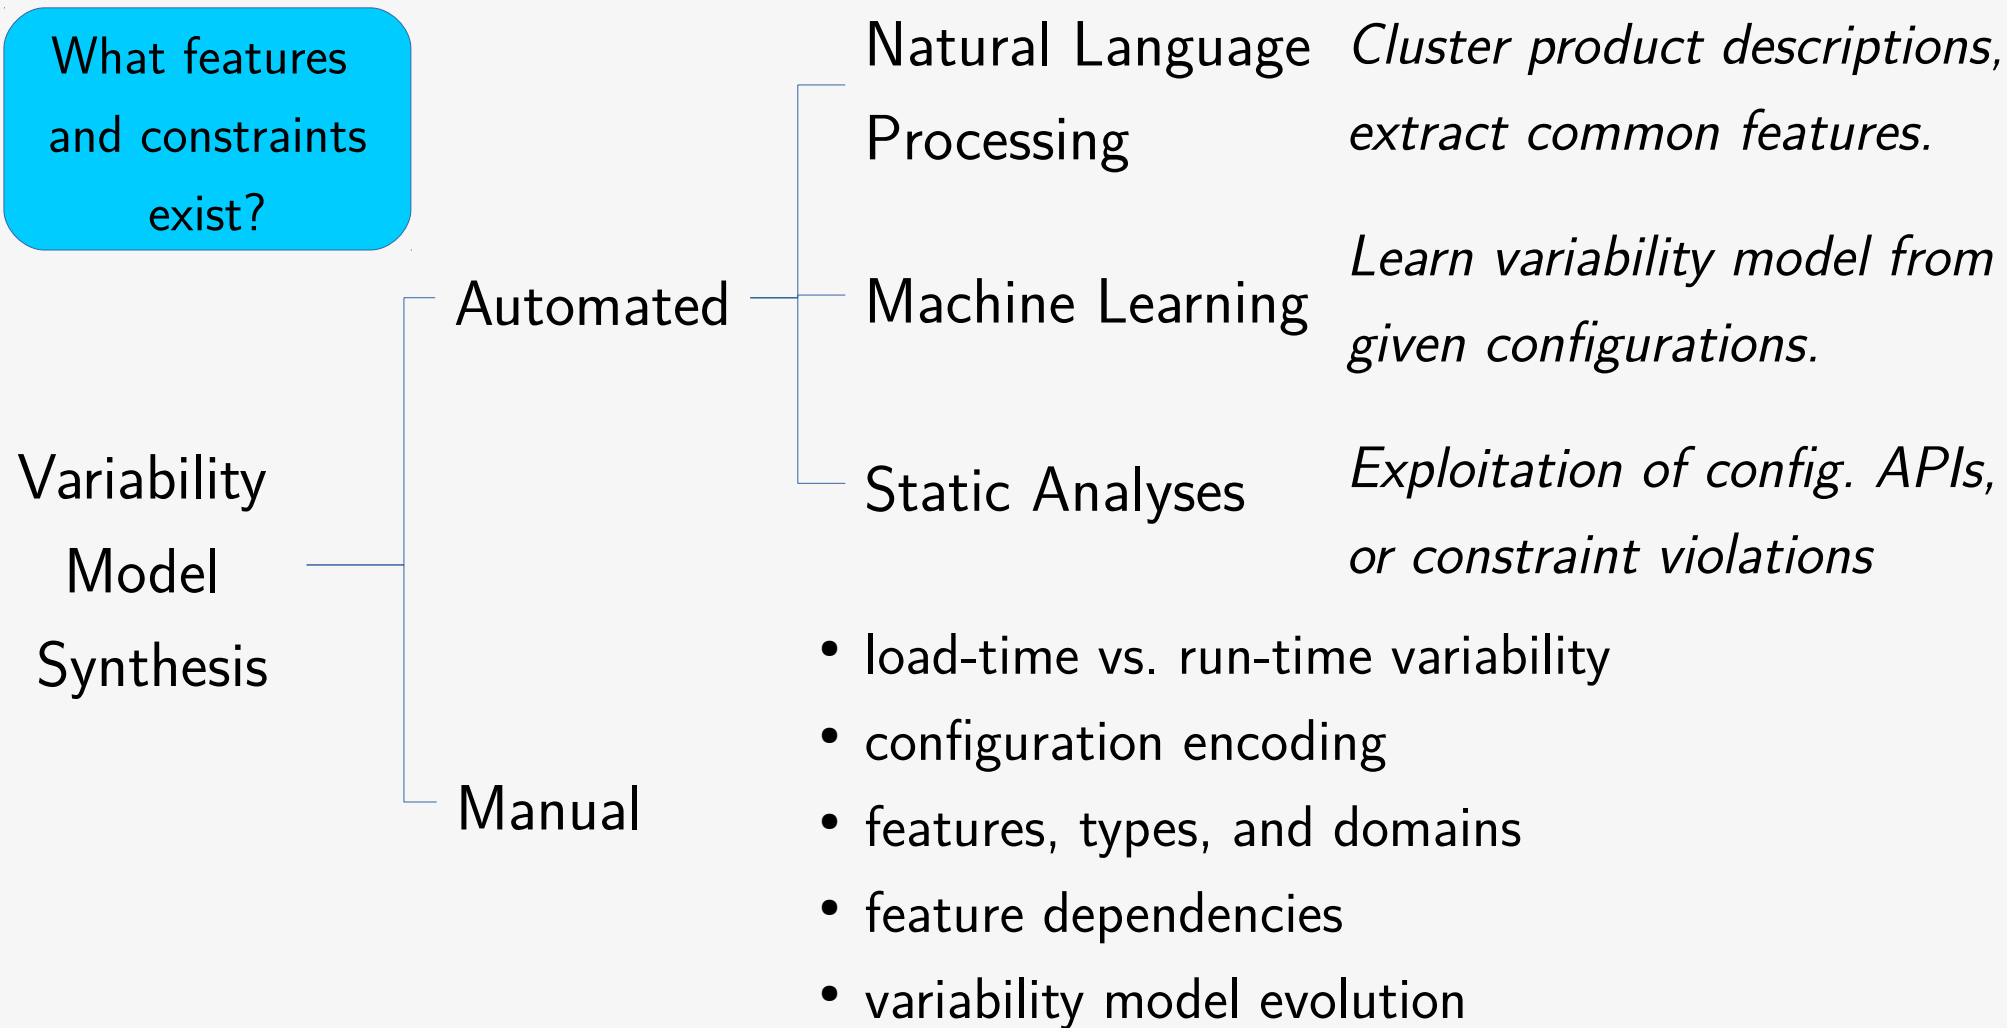
Problem: Performance Evolution Assessment

- Performance evolution: assessment of performance for multiple versions *required*
- Performance Evolution Assessment (PEA) – How to do that?
- Problem space is outlined by three intertwined dimensions:
 - Variability: infeasibly high numbers of variants, feature interactions
 - Performance Assessment: suitable performance indicators/measures
 - Diachrony: incremental code revisions, software development history


Goal: Methodology for PEA



Variability Model Synthesis (1)



Variability Model Synthesis (2)



- Automated approaches are only applicable under **preconditions**:
 - NLP: Textual description of a product required (\sim domain analysis)
 - ML: Sufficiently large number of valid feature selections required
 - Static analyses: highly-specialized use cases
- General purpose strategy: manual assessment based on documentation
- Configuration sampling with respect to feature interactions: t-wise sampling

Performance Assessment (1)



Performance
Benchmark

- Performance benchmark selection
 - Expressive: measure desired performance metrics/indicators
 - Reproducible: obtain similar results under equal circumstances
 - Cost-efficient: reasonable cost of benchmark testing
- Profiling: dynamic assessment of performance metrics
 - No general purpose profiler, depending on test setup
 - e.g., VisualVM for Java, AOP, network sniffers, ...
- Timing statistics on host-machine: `/usr/bin/time`

Performance
Testing Setup




Performance Assessment (2)



Robust/Applicable
Measurement
Summary

- Robust measures: median and interquartile range
 - Robust statistical measures are not distorted by extreme measurements
- Performance change magnitude:
Relative change per variant
- Performance change range:
variance of change per version

V1	0.5 %
...	...
V10	0.45 %
...	...
V20	-0.3 %




Evaluation & Case Study


RQ: What revisions induce significant performance changes?

- RQ: Does performance evolve for configurable software systems?
- GNU XZ: file compression tool
 - configurable at load-time, 9B/4N features
 - 36 variants sampled, ~ 1.100 versions assessed
- X264: video encoder
 - configurable at load-time, 8B/12N features
 - 8 variants sampled, ~ 2.800 versions assessed


Revision Assessment



- What revisions induce significant performance changes?
- Case study suggests 2 sampling strategies:
 - **Revision size**: more code revised → more likely performance changes
 - **Significant files**: certain files revised → likely impact on performance




Performance evolution: x264





- Effect magnitude: oscillations throughout all versions
 - small range of performance changes: - 0.015 % to 0.01 %
- Effect range: homogeneous evolution throughout all versions
 - X264 appears to be more mature than GNU XZ


GNU XZ: effect magnitude and range



X264: effect magnitude and range



Methodology overview



- Variability Assessment
 - automated vs. manual assessment
- Performance Assessment
 - performance metrics,
 - benchmarks,
 - profiler
 - robust measures, effect magnitude and range
- Revision Assessment
 - Sampling strategy: largest revisions or learn significant code units

Conclusion & Future Work

- Methodology description for configurable software systems
- Feasibility evaluation with a minor case study (2 systems)
- Insights obtained: performance evolution history
 - Possible indicator for software maturity and quality
- Use cases and further directions:
 - Performance prediction for future revisions and variants

Resources

- Articles

- <https://www.forbes.com/sites/ewanspence/2017/12/20/apple-iphone-kill-switch-ios-degrade-cripple-performance-battery/>
- <https://arstechnica.com/gadgets/2018/01/heres-how-and-why-the-spectre-and-meltdown-patches-will-hurt-performance/>
- <https://www.heise.de/newsticker/meldung/Intel-Benchmarks-zu-Meltdown-Spectre-Performance-sackt-um-bis-zu-10-Prozent-ab-SSD-I-O-deutlich-mehr-3938747.html>
- <http://bgr.com/2017/12/28/apple-batterygate-explainer-why-iphones-slow-down/>