

# Assessing Performance Evolution for Configurable Systems



Advisors: Prof. Dr.-Ing. Ina Schaefer, Prof. Dr.-Ing. Norbert Siegmund

September 20, 2017  
Stefan Mühlbauer

# Configurable Systems

- Software systems provide configuration options (features)
- (De-)selecting options tune, dis- or enable functionality



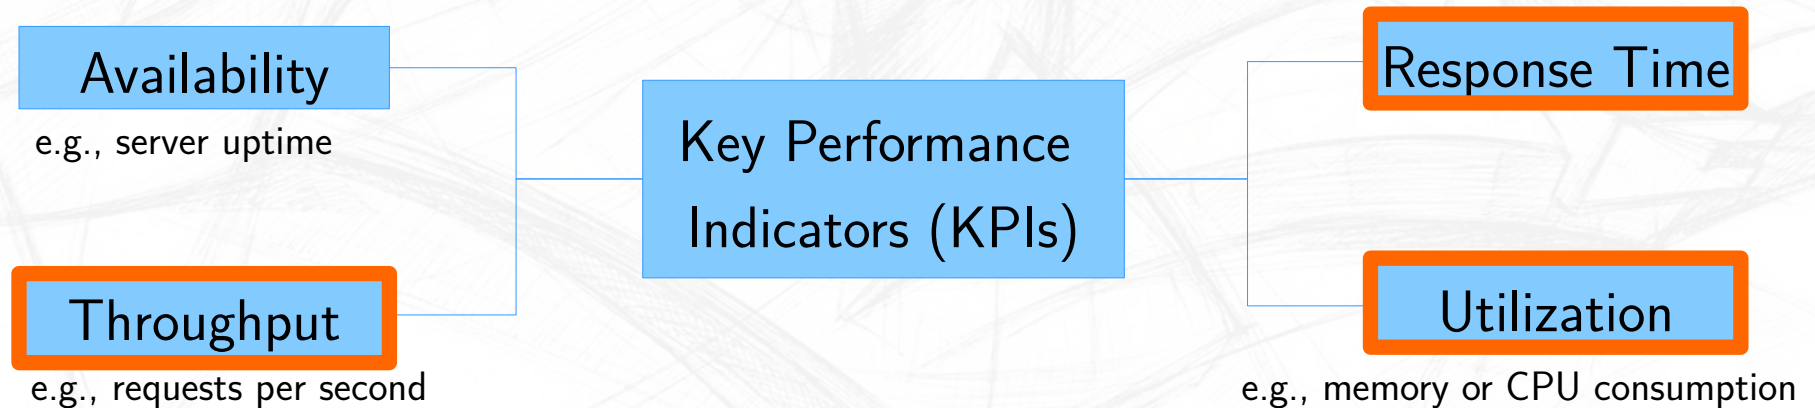
```
~$ zip -e -9 file.txt
```

- Unanticipated behavior can emerge with selections of multiple features (feature interaction)
  - Example: Compression and Encryption

*Compressing encrypted data can be **faster** than compressing raw data, since encrypted data is already more „compact“.*

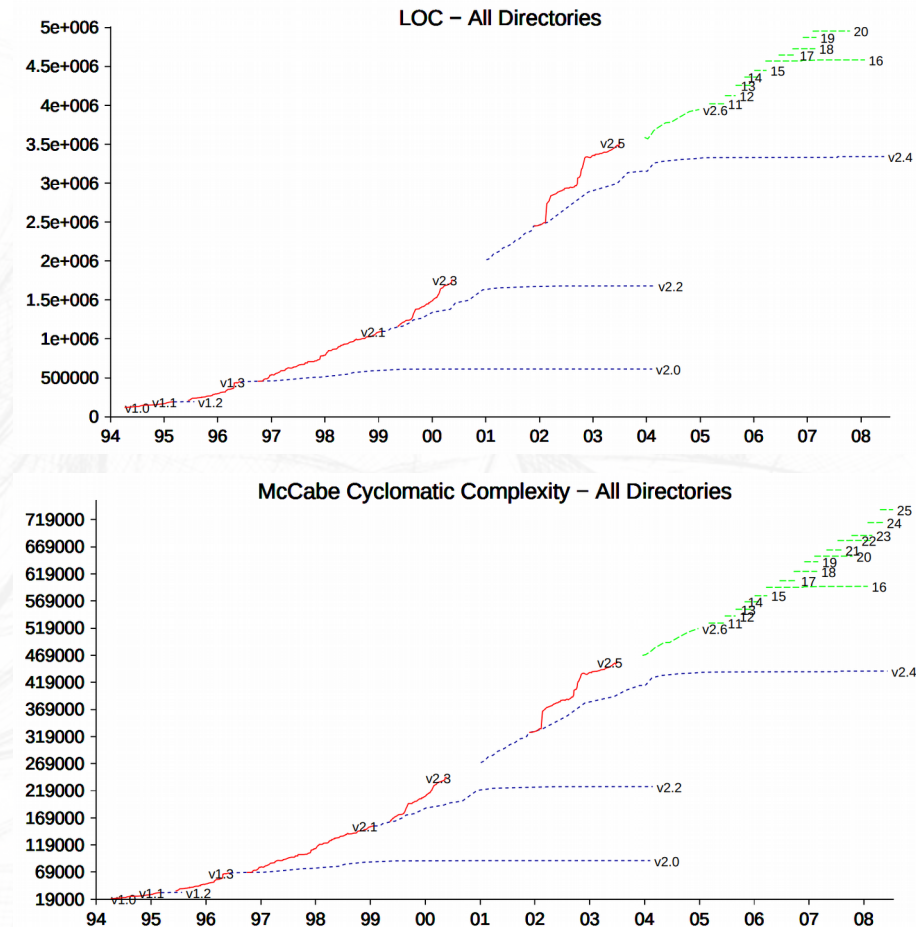
# Performance

- Performance: How successfully is a task being performed?
- Software performance is described by four categories:



# Software and Performance Evolution

- Evolution: Adaption to changing contexts/requirements
- Example: The Linux Kernel
- Performance regression
  - A symptom of software evolution
  - Degradation of performance quality of software over time



Source: Fig. 2 and 4 of [5]



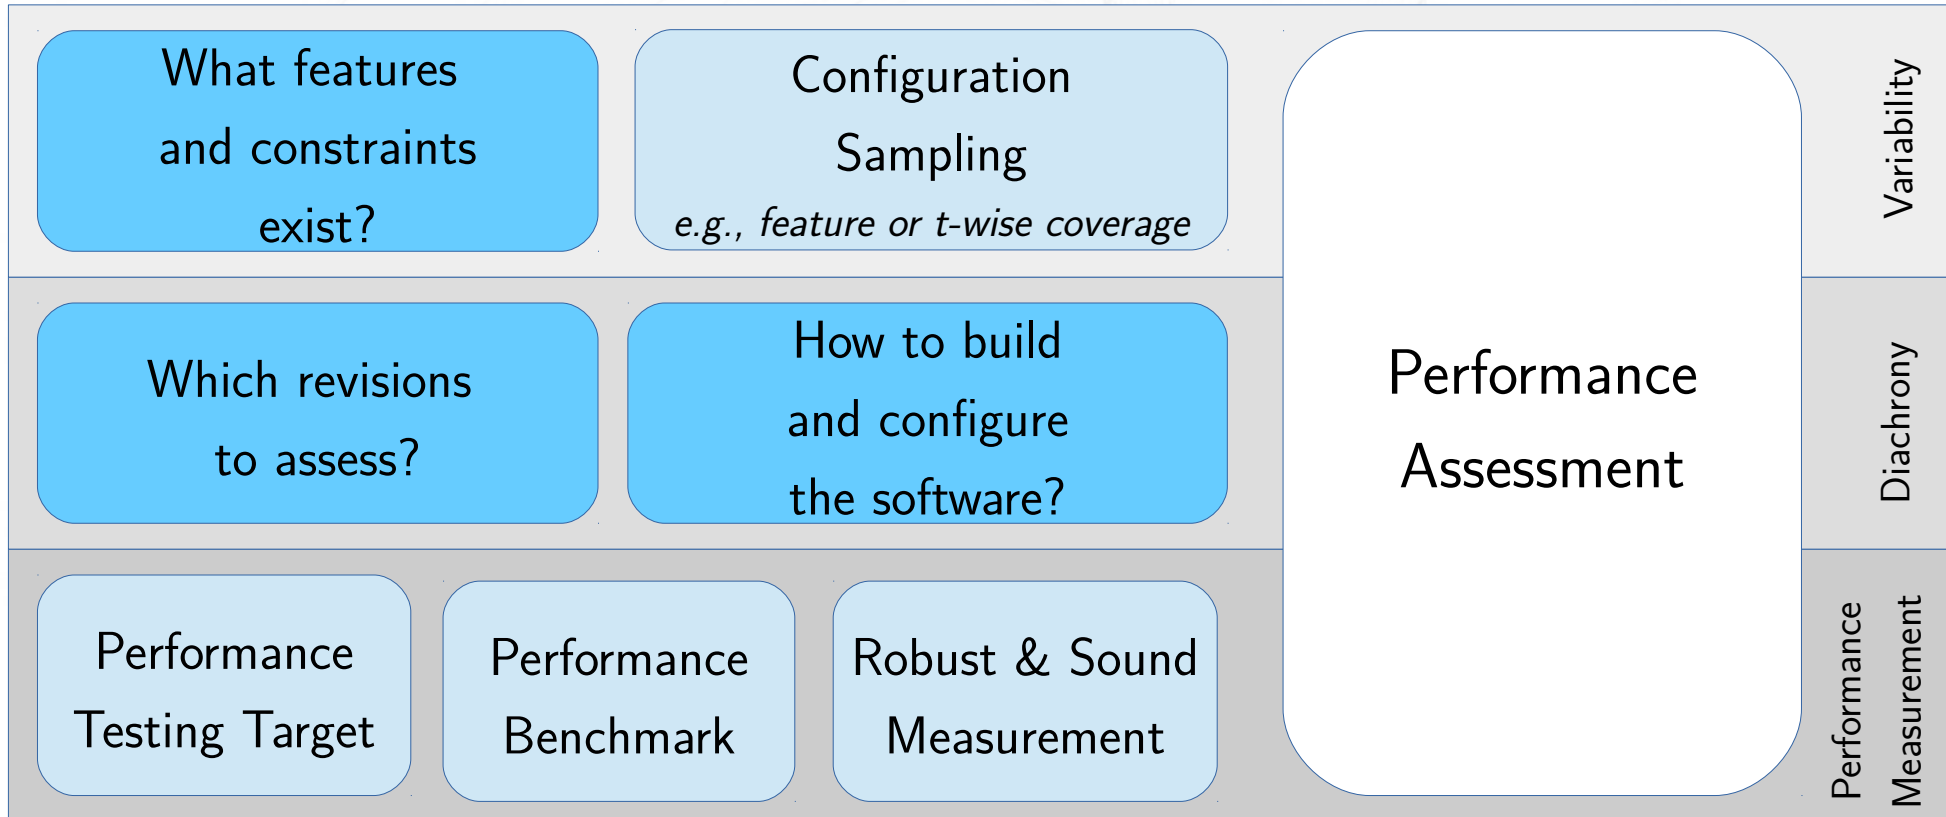
# Performance Assessment

- Performance evolution is the change of performance over time
- Performance assessment for multiple revisions **required** to capture performance evolution

# Motivation

- Problem: Performance Assessment – How?
- Goal: Description of performance assessment process with regard to variability, evolution, and statistical accuracy.
- Objectives
  - How to derive variability models and select configurations to assess?
  - How to select revisions to assess from revision history?
  - How to select robust and sound statistical measures?



# Performance Assessment Process



# Variability: Feature Model Synthesis



- Feature Extraction [2]
  - Exploit used configuration APIs
  - Recover feature names, types and domains
- Constraint Extraction [3]
  - Infer constraints from rule violations:
    - “Every valid configuration compiles”
    - “Every valid configuration yields a different product”
- *Optional*: Semi-automatically organize features and constraints to feature diagram [4]



$B \rightarrow A \vee \neg D$

$D \rightarrow C \vee \neg E$

$E \rightarrow C \vee \neg D$

$C \rightarrow A$

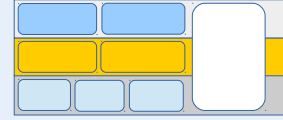


# Variability: Configuration Sampling

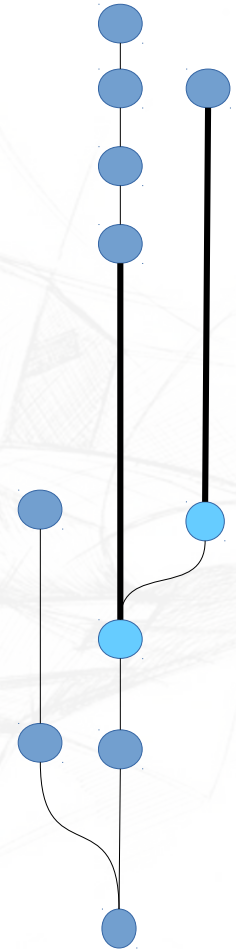


- Sampling: Finding a 'representative' subset of configurations
- Sampling strategies: Isolated solutions with coverage criteria
  - Pair- or t-wise: Coverage of simple and higher feature interactions
  - Feature coverage: Every feature selected at least once
  - ...

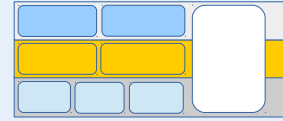
# Diachrony: Revision Sampling (1)



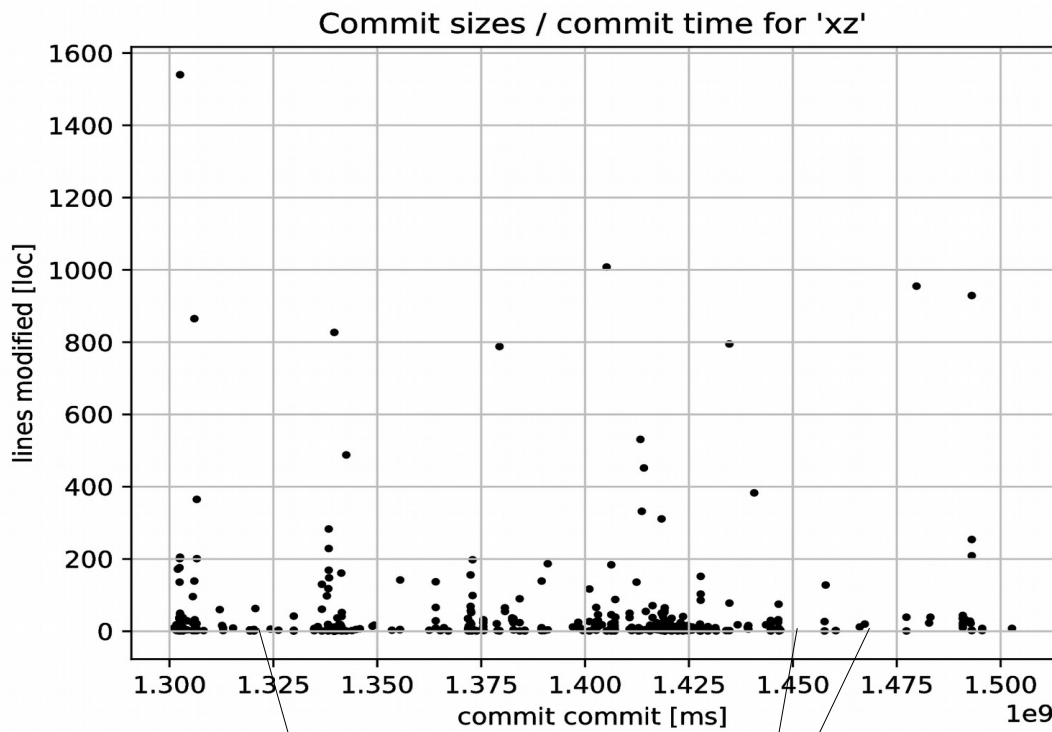
- Which revisions do we want to assess performance for?
  - Releases, release candidates,
  - bugs, and corresponding bug-fixes
- Information we can use to classify include
  - Revision history metadata (commit messages, ...)
  - Release notes
  - ...
- Which revisions best describe performance evolution?



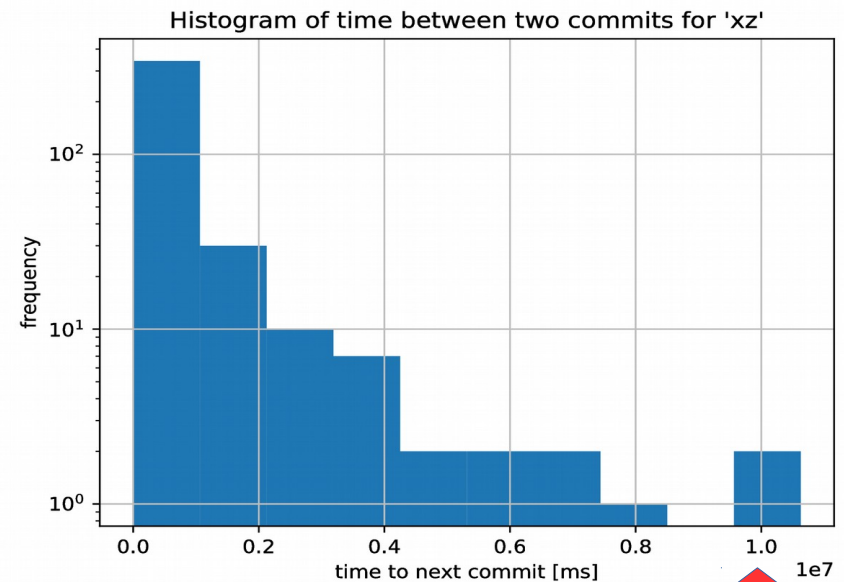
# Diachrony: Revision Sampling (2)



- Idea: Revisions which have been the latest ones for a long time



„commit streaks“ followed by sparse sections



Very few commits with high temporal distance to next commit




# Semi-automated Integration



- How to automatically build a software system?
  - Usually: Manually predefined build routine, e.g., shell script
  - Possible extensions: Pattern matching for 'characteristic' files?
- How are configurations read by the software system?
  - Usually: Manually predefined templates, e.g., .properties file
  - Possible extensions: Pattern matching for 'characteristic' files?



# Measurement: Soundness



- How to describe performance statistically accurately?
- Soundness: Are conclusions drawn from measurements correct?
  - Example: Use harmonic mean for rates, not the arithmetic mean
    - Two machines processing 100 requests with 100 and 10 hits/s each
    - Arithmetic mean: 55 hits/s, harmonic mean: 18.182 hits/s
    - Time required for processing requests:  $1\text{ s} + 10\text{ s} = 11\text{ s}$
    - Two 'average' machines should take 11 seconds
    - Arithmetic mean: 3.64 s, harmonic mean: 11 s

# Measurement: Robustness



- Robustness: Is the measure affected by outliers?
  - Example: Series 1, 1, 2, 3, 8, arithmetic mean of 5, median of 2
- Robust measure of central tendency: Median (2<sup>nd</sup> quartile)
  - Bigger than the lower 50% of measurements
  - Smaller than the upper 50% of measurements
- Robust measure of dispersion: Inter quartile range (IQR)
  - Difference between the 3<sup>rd</sup> and 1<sup>st</sup> quartile

# Outlook for Evaluation



- Evaluation of the aforementioned classification problems
  - arbitrary systems selected from GitHub?
- Implementation of a performance assessment tool chain
- Modeling of performance behavior for 2 – 3 smaller systems
  - GNU xz utils (free file compression tool)
  - x264 (free audio and video encoder)
  - ...



# Literature

- [1] Molyneaux, I. (2014). *The Art of Application Performance Testing: Help for Programmers and Quality Assurance* (2nd ed.). O'Reilly Media, Inc.
- [2] Rabkin, A., & Katz, R. (2011). *Static extraction of program configuration options*. In Proceedings of the 33rd ICSE (pp. 131–140). ACM.
- [3] Nadi, S., Berger, T., Kästner, C., & Czarnecki, K. (2015). *Where do configuration constraints stem from? an extraction approach and an empirical study*. IEEE Transactions on Software Engineering, 41(8), 820–841.
- [4] She, S., Lotufo, R., Berger, T., Wasowski, A., & Czarnecki, K. (2011). *Reverse engineering feature models*. In Software Engineering (ICSE), 2011 33rd ICSE (pp. 461–470). IEEE.
- [5] Israeli, A., & Feitelson, D. G. (2010). *The Linux kernel as a case study in software evolution*. Journal of Systems and Software, 83(3), 485–501.