Technische Universität Braunschweig

Department of Computer Science

Master's Thesis

Assessing Performance Evolution Of Configurable Software Systems

Untersuchungen zur Performanz-Evolution von konfigurierbaren Software-Systemen

Author:

Stefan Mühlbauer

August 27, 2017

Advisors:

Prof. Dr.-Ing. Ina Schaefer

Technische Universität Braunschweig · Institute for Software Engineering and Automotive Informatics

Prof. Dr.-Ing. Norbert Siegmund

Bauhaus University Weimar · Chair for Intelligent Software Systems

Eidesstattliche Erklärung

Hiermit versichere ich eidesstattlich, dass ich die Arbeit selbständig verfasst und keine anderen als die angegebenen Quellen und Hilfsmittel benutzt habe.

Braunschweig, August 27, 2017   _____

Statement of Originality

This thesis has been performed independently with the support of my supervisors. To the best of my knowledge, this thesis contains no material previously published or written by another person except where due reference is made in the text.

Braunschweig, August 27, 2017   _____

Contents

1. Introduction

Configurable Software   Modern software systems often need to be customized to satisfy user requirements. Customizable software, for instance, enables greater flexibility in supporting varying hardware platforms or tweaking system performance. To make software systems configurable and customizable, they exhibit a variety of *configuration options*, also called *features* (Apel et al., 2013). Configuration options range from fine-grained options that tune small functional- and non-functional properties to those that enable or disable entire parts of the software system. The selection of configuration options can be accommodated at different stages, either at *compile-* or *build-time* when the software is built or at *load-time* before the software is actually used. Compile-time variability usually governs what code sections get compiled in the program. For instance, compile-time variability can be realized by excluding code sections from compilation using preprocessor annotations (Hunsen et al., 2016) or by assembling the code sections to compile incrementally from predefined code modules depending on the feature selection (Schaefer et al., 2010). In contrast to that, load-time variability controls which code sections can be visited during execution. Configurations for load-time variability can be specified using configuration files, environment variables or command-line arguments. Many software systems are configurable, examples range from small open-source command-line tools to mature ecosystems including Eclipse or even operating systems such as the Linux kernel with more than 11,000 options (Dietrich et al., 2012).

Configuration options for software systems are usually constrained (e.g., are mutually exclusive, imply or depend on other features) to a certain extent. In the worst case though, where all options can be selected independently, the number of valid configurations grows exponentially with every feature added, and exceeds the number of atoms in the entire universe once we count 265 independent features. Hence, even for a small number of features, any naive approach for assessing emergent properties of configurable software systems exhaustively for each valid configuration is conceived infeasible. Despite this mathematical limitation, many feasible approaches to static analysis for configurable systems emerged. Those variability-aware approaches enable, for instance, type checking in the presence of variability by exploiting commonalities among different variants (Thüm et al., 2014).

To meet functional and non-functional requirements, users aim at finding the optimal configuration of a configurable software system. However, this task is non-trivial and has shown to be NP-hard (White et al., 2009). The main driver for the complexity are feature interactions. A *feature interaction* is an "emergent behavior that cannot be easily deduced from the behaviors associated with the individual features involved" (Apel et al., 2013) and can make development and maintenance of a configurable system an error-prone task.

To illustrate feature interactions, consider the following example (Siegmund et al.,

2015). A software system, say a file server, is used to store files in a database and provide access upon request. The system provides two features, encryption and compression. In isolation, both file en- or decryption and file (de-)compression demand an expectable fraction of memory and processor time. The performance behavior for the software system though may vary if both features are selected. For instance, if a file is encrypted and compressed (or vice versa), we can expect the operation to demand less resources since an encrypted file is likely to be of smaller size than the decrypted original. This is a positive example for a feature interaction, where the performance behavior although being benefitting, is unexpected.

Performance Behavior   The term "performance" with respect to software and software systems is not precisely defined and differs from an end user's and a developer's perspective. According to Molyneaux (2014), from a user's perspective "a well-performing application is one that lets the end user carry out a given task without any undue perceived delay or irritation". However, to accurately assess performance, from a practitioner's perspective, performance is outlined by measurements called *key performance indicators* (KPIs) which relate to non-functional requirements (Molyneaux, 2014). The set of KPIs includes availability of a software system, its response time, throughput, and resource utilization. Availability comprises the amount of time an application is available to the user. Response time describes the amount of time it takes to process a task. Throughput describes the program load or number of items passed to a process. Resource utilization describes the used quota of resources required for processing a task.

The performance behavior of a software system depends on the functionality offered, the respective implementation, program load, the underlying hardware system, environment variables, and the resulting execution. Since configuration options control what and how functionality is executed, we concentrate here on those source of performance. While feature interactions not necessarily cause the software system to break severely in all cases, its overall performance can become unfavorable for corner cases or specific configurations as the feature selection influences the execution. That is, the choice of features as well shapes the performance of a software system.

Performance And Evolving Software   Actively maintained software systems evolve with every modification made, every version released, and patch provided. Modifications usually introduce new functionality to the system, but functionality might as well be divided into smaller modules to enable provide more fine-grained configuration options. When features are removed from the software system, the corresponding functionality might remain in the code base or options are merged (Apel et al., 2013).

There exists substantial work on understanding the evolution of configurable systems, for instance, with respect to software architecture (Zhang et al., 2013; Passos et al., 2015) or variability (Seidl et al., 2012; Peng et al., 2011; Passos et al., 2012). As software evolves, the code base which is subject to modifications and the overall architectural quality can degrade. Common symptoms of architectural degradation are code tangling and scattering (Zhang et al., 2013; Passos et al., 2015), which lead to less cohesive and stricter coupled code. The more the code base is constrained and interdependent, the more software can become "brittle" (Perry and Wolf, 1991), less flexible, harder to adapt, and therefore harder to evolve. Evolution of software, especially with respect to variability, is essentially driven by and can be conceived as adapting a software system to changed

requirements and contextual changes (Peng et al., 2011). That is, (potential) degradation of software quality as software evolves is often a phenomenon due to decisions trading quality assurance (QA) and maintenance tasks with meeting requirements and schedules (Guo et al., 2011). The metaphor of *technical debt* (Guo et al., 2011) which is commonly used to describe this trade-offs and corresponding costs, outlines the risk that postponed maintenance tasks pose to software evolution. Although every deferred maintenance or QA task may save some cost, it also could have unveiled software defects in the first place. Technical debt implies both interest, so to speak, the potential damage of a defect depending on its severity, as well as the probability of incurring interest. A defect can be severe, yet fixable with reasonable effort and cost. However, the aforementioned symptoms of architectural degradation and deferring maintenance render bug-fixing to become more and more expensive.

Besides the aspects of software evolution discussed above, the evolution of performance for software systems has gained more attention recently. In practice though, quality assurance with respect to performance is still conducted to an unsatisfactory extent, or accommodated to late in the development process, according to Molyneaux (2014). Thus, postponed maintenance and QA with respect to performance is likely to a driving factor for degradation of performance quality, or simply called *performance regression*.

While performance discipline has emerged as a discipline of or target in software testing, qualitative root cause analysis, for the most part, is still conducted manually (Molyneaux, 2014). However, there exists work on automated root cause analysis for performance bugs, such as measuring the execution time of unit tests, whereby an increased execution time indicates performance regression, and the corresponding unit test helps isolating the root cause thereof (Heger et al., 2013; Nguyen et al., 2014). In conclusion, we see that, to better assure good software performance, more knowledge about performance behavior needs to be available, ideally, earlier in the development process.

Performance Prediction For configurable systems, performance behavior can be more complex and dependent on the feature selection, as we have seen with the example for feature interactions above. Similarly, quality assurance for configurable software systems is far from exhaustively testing all possible configurations, but rather close to only testing a selection of configurations sampled with respect to certain constraints. Sampling strategies might stress feature interactions, such as pair-wise sampling (Apel et al., 2013). However, all sample are selected with the intention to learn as much as possible about the entire system from a small sample of variants. So to speak, a sampling strategy is "optimal" if for a resulting sample, the probability of missing an arbitrary (relevant) feature interaction, is minimal.

While performance testing is apparently useful, recently, a number of techniques to model and predict performance behavior for arbitrary configurations have emerged. The idea behind all these approaches is similar to sampling strategies as an "accurate" prediction model is able to predict performance for arbitrary configurations with a low error rate. The underlying optimization problem of performance prediction models is to find an accurate estimator $\tilde{f}$ for a function f describing a performance property depending on the feature selection. Performance properties can be estimated without performance measurements, for instance by inferring performance properties from software models (Woodside et al., 2007), measurement-based approaches for configurable systems address this optimization problem. The proposed approaches include performance prediction models

build on learning performance behavior with decision trees (Guo et al., 2013), learning a frequency-based representation of the target function (Zhang et al., 2015), or learning the influence of single features and all performance-relevant feature interactions (Siegmund et al., 2012, 2015). All approaches have shown promising error rates for several real-world applications and allows prediction of system performance for arbitrary configuration variants. However, all approaches to create performance prediction models demand samples of performance measurements for multiple configurations to learn performance behavior and validate predictions.

Problem Statement The assessment of performance evolution requires a series of performance prediction models describing performance behavior for a series of versions of a configurable system. Assessing the performance behavior for a single version of a configurable software system entails a number of necessary and preliminary tasks. These tasks can even become more complicated once, instead of a single version, a series of versions is assessed:

- □ *Feature Model Synthesis:* Not all configurable systems do explicitly exhibit a variability model what is required to derive all valid variants (Rabkin and Katz, 2011; Nadi et al., 2015). While substantial work exists on reverse engineering variability models from code (Rabkin and Katz, 2011; She et al., 2011; Zhou et al., 2015; Nadi et al., 2015) or non-code artifacts (Alves et al., 2008; Andersen et al., 2012; Bakar et al., 2015), many techniques still involve manual decisions (She et al., 2011) and domain knowledge (Nadi et al., 2015). Moreover, variability models evolve as part of the software (Peng et al., 2011), vary from version to version, and therefore, require repeated reverse engineering steps.

- □ *Configuration Translation:* the translation of a valid configuration to a configuration artifact such as a configuration file or a list of command-line arguments may differ from system to system. This step may be automated, but one still needs to detect how configurations are read by the software system one wants to study.

- □ *Automated Integration:* Same goes for the infrastructure to compile or build a software system since there exist may possible build tools such as makefiles or sbt. Again, the build process can be automated, but one needs to detect and specify the build mechanism used.

- □ *Version History Sampling:* To study performance evolution one needs to specify which snapshots or versions of a software system one wants to study. While detecting releases and release candidates should be straightforward, one might, for instance, be interested in the performance evolution including snapshots between two releases. As not all snapshots though are likely to compile, classifying defect snapshots can still be tedious work.

- □ *Performance Assessment Setup:* The accurate assessment of performance evolution requires a suitable testing setup. The methodology required for assessing performance among others requires the selection of suitable performance metrics and corresponding benchmarks, means to record measurements and repeat experiments easily, and proper ways to interpret and compare results.

Goals And Thesis Structure   The goal of this thesis is to provide a theoretical and practical foundation for exhaustive performance measurements of configurable software systems and series thereof. We contribute a guideline of and tool support for performance measurements for configurable and evolving software systems. Our research objectives and desired outcomes are

- a literature overview regarding software evolution, feature model synthesis and performance assessment,

- a methodology to assess performance evolution with respect to the aforementioned challenges, and

- a practical tool for performance measurement for multiple revisions of configurable software systems.

The Thesis is organized as follows. Chapter 2 provides the background to the relevant topics discussed in this thesis, including variability modeling, software evolution, feature model synthesis, performance assessment, statistical basics to summarize data records, and performance prediction models. In Chapter 3, we propose our measurement methodology and discuss the methods used for our performance measurement tool as well as its limitations. In Chapter 4, we evaluate several aspects of our tool with respect to practicality and discuss the results thereof. Finally, Chapter 5 concludes the thesis and gives an outlook on possible future work.

2. Background

This chapter is intended to recapitulate the background of the thesis theme. In Sec. 2.2, we recall the evolution of software systems with respect to architecture and variability. In Sec. **??** we outline the characteristics of software performance, practical testing and measurement strategies as well as some statistical background necessary to analyze, interpret and compare performance assessment. In Sec. 2.3 we present recent approaches for feature model extraction from existing software systems and code artifacts. Finally, in Sec. **??** we recall and compare in detail different approaches to model and predict performance behavior for configurable software systems.

2.1 Variability Modeling

The design and development of configurable software systems is conceptually divided into *problem space* and *solution space* (Czarnecki and Eisenecker, 2000). The problem space comprises the abstract design of features that are contained in the software system as well as constraints defined among features, such as dependencies or mutual-exclusion. The solution space describes the technical realization of features and the functionality described by and associated with features, e.g., implementation and build mechanisms. That is, features cross both spaces since they are mapped to corresponding code artifacts.

A common way to express features and constraints in the problem space is to define a *variability model*, or *feature model*, which subsumes all valid configurations (Kang et al., 1990; Apel et al., 2013). There are different and equivalent syntactical approaches to define feature models, for instance, a propositional formula F over the set of features of the configurable software systems (Batory, 2005). In this case a configuration is valid with respect to the feature model if and only if F holds for all selected features being true and all unselected features being false respectively. However, a more practical and more commonly used way to express feature models are graphical tree-like *feature diagrams* (Apel et al., 2013). In a feature diagram, features are ordered hierarchically, starting with a root feature and subsequent child features. By definition, the selection of a child feature requires the parent feature to be selected as well. Child features can either be labeled as *optional* features or *mandatory* features; the latter ones need to be selected in every configuration. Moreover, feature diagrams provide a syntax for two different types of feature groups, *or-groups* and *alternative-groups*. For an or-group, at least one of the group's features needs to be selected for a valid configuration, whereas for an alternative group exactly one out of the group's mutually exclusive features must be selected. In addition to the feature hierarchy, constraints, which cannot be expressed by the tree-like structure, are referred to as *cross-tree constraints*. Cross-tree constraints, depending on the notation, are depicted by arrows between two features or simply added to the feature

diagram as a propositional formula. For such two features f_1 and f_2, a cross-tree constraint means that for feature f_1 to be selected, either the selection of f_2 is required/implied or excluded.

An introductory example for the syntax and semantics of feature diagrams is provided in Fig. 2.1. In this example an imaginary vehicle propulsion can be configured with eight valid configurations. The vehicle requires an engine, and therefore feature Engine is mandatory. At least one out of the three features Hybrid, Piston and Electric needs to be selected. For a piston engine, we can select either the feature Diesel or Petrol. A petrol engine requires additional ignition sparks in contrast to a Diesel engine. For an electric engine we require a battery, hence, the feature Battery is mandatory. In addition, the feature model specifies two cross-tree constraints: First, the feature Tank is optional, yet once a piston engine is selected, we require a tank. Second, if we want to use the Hybrid functionality (e.g., use both electric and piston engine simultaneously), we require to have both a piston and an electric engine.

Hybrid ⇒ Electric ∧ Piston

Piston ⇒ Tank
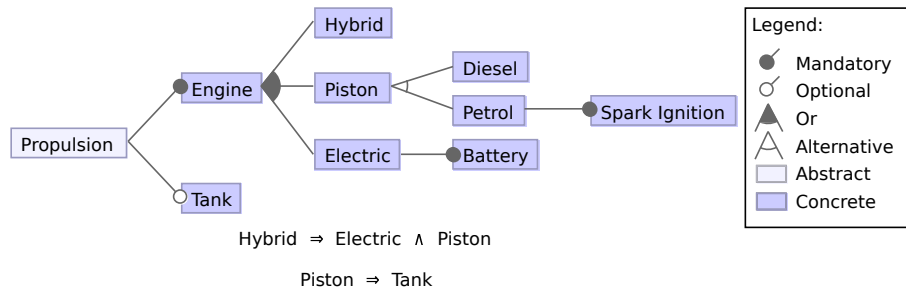
Fig. 2.1. Feature diagram for a feature model with eight valid configurations; two cross-tree constraints are specified as propositional formulas over features

2.2 Evolving Software

2.2.1 Software Evolution

The first notion of the software development process is usually developer-centered and merely focuses on software being designed, implemented, tested, and eventually being released and deployed. Maintainability is a generally recognized software quality property to look after, and maintenance is, of course, essential to every successful software system. Nonetheless, less attention is given to the ability to adapt a software system to changing requirements (evolvability) rather than maintaining it to keep functionality working (Parnas, 1994). Software evolution and evolvability, like software itself are manifold. Software evolves in many ways ranging from maintenance (refactoring, bug-fixes and patches) to adapting to changed requirements (adding, removing, reorganizing functionality and variability). Modern software systems not only often ship with a variety of configuration options to select from, they also employ routines to be build and sometimes even make use of or are part of platforms, such as apps or plugins. That is, software evolution affects all aforementioned aspects, maintainability as well as evolvability can degrade as software evolves.

2.2.2   Software Erosion

The negative symptoms of software evolution which are referred to as "architectural erosion" (Breivold et al., 2012) have been addressed by many researchers. Most of existing research so far though focuses on evolution regarding software architecture (Breivold et al., 2012). The main driving factors leading to symptoms of decay identified by Perry and Wolf (1991) are architectural erosion and architectural drift. While architectural drift subsumes developers' insensitivity when not following a systems architecture or respective guidelines while making changes, architectural erosion subsumes ignoring and violating the existing software architecture. Parnas (1994) argues that as software evolves, software is maintained and evolved by developers who are not necessarily familiar with the initial architectural design. Therefore, knowledge about the architecture can become unavailable. Although the unfavorable effects of software evolution do not necessary break a system necessarily and imminently, the software becomes "brittle" (Perry and Wolf, 1991) as maintainability as well as evolvability degrade. Concrete symptoms of software erosion on the implementation level have been documented.

Zhang et al. (2013) have studied erosion symptoms for a large-scale industrial software product line with compile-time variability using preprocessor directives. The authors identify variability-related directives and clusters of those that tend to become more complex as the software evolves. The negative effects, or symptoms of software erosion are described as, but not limited to *code replication* and interdependencies between code elements, such as *scattering* and *tangling*. Code scattering describes the phenomenon of code belonging to a certain feature being scattered across multiple units of implementation, e.g., modules, whereas code tangling means that code from different and potentially unrelated features in entangled within a single module.

Passos et al. (2015) have studied the extent of usage of scattering for device-drivers in the Linux kernel. Despite scattering being quite prevalent, their findings suggest that the kernel architecture is robust enough to have evolved successfully. Nonetheless, platform drivers in the Linux kernel seem more likely to be scattered than non-platform driver. They conclude that this is a trade-off between maintainability and performance: a more generalized and abstract implementation for platform-drivers in this case could possibly avoid scattering, yet refactorings in this manner did not seem to be necessary or worth the effort yet.

2.2.3   Variability Evolution

Apart from architecture evolution, the variability offered by software systems evolves as well. For configurable software systems, evolution steps will not only affect artifacts in the solution space, yet also be visible in changes in the respective variability models in the problem space. Although the variability aspect of software evolution has not gathered as much attention as architecture in the past, more and more research has emerged recently to address and understand variability evolution.

Peng et al. (2011) proposed a classification of variability evolution patterns that conceives evolution as adaption to changing (non-)functional requirements and contexts. For a context in that sense, two categories exist. A driving context determines, whether a variability model and respective variants can meet functional requirements in the first place. A supporting context by definition determines how non-functional properties are strengthened or weakened. Any changed requirement is likely to change the contexts for a

software systems variability model and, therefore, will make adaptations of the variability model necessary. Within their classification method, Peng et al. (2011) identify major causes for variability evolution, comprising a) new driving contexts emerging, b) weakened supporting contexts (for instance, due to new non-functional requirements), and c) unfavorable trade-offs for non-functional properties.

To understand single evolutionary steps, several catalogs of variability evolution patterns have been proposed. Peng et al. (2011) present three patterns, where either a new feature is added, a mandatory feature becomes optional, or a mandatory/optional feature is split into alternative features. Seidl et al. (2012) suggest a catalog of patterns for co-evolution of variability models and feature mappings that additionally introduces code clones, splitting a feature into more fine-grained sub-features and feature removal as evolution patterns. In addition, Passos et al. (2012) have studied variability evolution in the Linux kernel and present a catalog of patterns where features are removed from the variability model, but remain a part of the implementation. Their catalog, among others, includes feature merges, either implicit (optional feature merged with its parent) or explicit.

The classification proposed by Peng et al. (2011) is a general and formalized approach that, as well as Seidl et al. (2012) and Passos et al. (2012), describes elementary evolution patterns which can be composed to more complex patterns. Nonetheless, no comprehensive catalog of variability evolution so far has been proposed.

2.3 Feature Model Synthesis

A variability model as an abstraction of functionality of a software system is required, or at least of great interest, in many contexts.

First, not every configurable system provides an explicit representation of its variability model. The reasons for inexplicit or absent configuration specification are manifold. They can range from poor or inconsistent documentation (Rabkin and Katz, 2011) to overly complex configurability (Xu et al., 2015), or configuration constraints originated in different layers of a software system, e.g.m build constraints or compiler constraints (Nadi et al., 2015).

Second, variability models have emerged to be a useful means in domain analysis prior to developing a software system. As variability models group and organize functionality, synthesizing a variability model has shown to be applicable to extract features and constraints from functional requirements. In addition, by comparison of product specifications for an existing market domain, variability models can provide detailed feature summary.

For this thesis, we focus on the first aspect of synthesizing variability models as our work addresses the assessment of already existing configurable software systems. Nonetheless, many techniques employed in the aforementioned second aspect address similar problems, yet rely on natural language artifacts rather than code artifacts (Alves et al., 2008; Bakar et al., 2015). The following section recalls work on extracting configuration options and constraints from source code as well as the organization of constraints into feature hierarchy and groups. The further assessment of configurable systems requires a well-defined and sound variability model.

2.3.1   Feature Extraction

The first objective in recovering a variability model from a configurable system is to determine the available configuration options to select from. In addition, for further configuration, the type of each configuration option (e.g., boolean, numeric or string) and the respective domain of valid values needs to be specified.

Rabkin and Katz (2011) proposed a static, yet heuristic approach to extract configuration options along with respective types and domains. Their approach exploits the usage of configuration APIs and works in two stages. It commences with extracting all code sections where configuration options are parsed. Next, configuration names can be recovered as they are either already specified at compile-time or can be reconstructed using string analysis yielding respective regular expressions. Moreover, the authors employ a number of heuristics to infer the type of parsed configurations as well as respective domains. First, the return type of the parsing method is likely to indicate the type of the configuration option read. Second, if a string is read initially, the library method it is passed to can reveal valuable information about the actual type. For instance, a method *parseInteger* is likely to parse an integer value. Third, whenever a parsed configuration option is compared against a constant, expression, or value of an enum class, these might indicate valid values or at least corner cases of the configuration option domain. The extraction method by Rabkin and Katz (2011) is precise, but limited, for instance, when an option leaves the scope of the source code. Nonetheless, for the systems studied, the authors recovered configuration options that were not documented, only used for debugging or even not used at all.

2.3.2   Constraint Extraction

The second step in recovering a variability model is the extraction of configuration constraints. An approach proposed by Zhou et al. (2015) focuses on the extraction of file presence conditions from build files using symbolic execution. A more comprehensive investigation of configuration constraints and their origin is provided by Nadi et al. (2014, 2015). They propose an approach based on variability-aware parsing and infer constraints by evaluating make files and analyzing preprocessor directives. Inferred constraints result from violations of two assumed rules, where a) every valid configuration must not contain build-time errors and b) every valid configuration should result in a lexically different program, thus. While the first rule aims at inferring constraints that prevent build-time errors, the second one is intended to detect features without any effect, at least as part of some configurations. Their analysis one the one hand emerged to be accurate in recovering constraints with 93 % for constraints inferred by the first rule and 77 % for second one respectively. On the other hand, their approach recovered only 28 % of all constraints present in the software system. Further qualitative investigation, including developer interviews, lead to the conclusion that most of existing constraints stem from domain knowledge (Nadi et al., 2015).

2.3.3   Feature Hierarchy Recovery

Besides recovering configuration options and respective constraints, to reverse engineer a feature model, one further step is required. The recovered knowledge needs to be organized in a tree-like hierarchy with feature groups specified and CTCs explicitly stated to derive a

valid feature diagram (Kang et al., 1990). While several approaches to the recover feature model hierarchy have been proposed, we are primarily interested in finding a hierarchy for knowledge obtained from source code. Other scenarios, as already stated in the opener of this section, are based on product descriptions or sets of valid configurations. The remainder of this subsection we will focus on organizing features and constraints extracted from source code. For further reading, Andersen et al. (2012) present algorithms for structuring feature diagrams for three different scenarios including the ones previously mentioned.

Given an extracted set of features along with corresponding descriptions and recovered constraints among the features, She et al. (2011) propose a semi-automated and interactive approach to synthesize a feature hierarchy. Their approach comprises three tasks. First, an overall feature hierarchy based on feature implications is specified. Second, potential feature groups are detected and manually selected. Finally, the feature diagram is extended with remaining CTCs. A detailed decription of the algorithm is presented below:

1. The algorithm commences with finding a single parent for each feature and, thus, specifying a tree-like feature hierarchy. Based on the given constraints a directed acyclic graph (DAG) representing implication relationships among features, a so-called implication graph, is constructed. Every vertex in the implication graph represents a feature and edges are inserted for each pair of features (u, v), where $u \implies v$ holds with respect to the given constraints.

   In addition to the implication graph, the algorithm for each feature computes two rankings of features that are likely to be the respective parent feature. The two rankings both employ the feature descriptions. Feature descriptions are compared for similarity using a similarity metric. For two features p and s the similarity is defined as the weighted sum of the inverse document frequencies $idf(w)$ for the words that both descriptions of features p and s share. The idf-ranking for a word w is the logarithm of the number of features divided by the number of features whose description contains w. Each idf value is weighted with the frequency of w in the description of feature p.

   The first ranking, called Ranked-Implied-Features (RIF), for each feature f ranks all features by their similarity to f in an descending order, but prioritizes those features that are implied according to the previously computed implication graph. The second ranking, called Ranked-All-Features (RAF) is similar to RIF, yet less strict since implied features are not prioritized. Given these rankings, a user for each feature selects a suitable parent feature from the RIF or RAF ranking. The idea behind providing two separate rankings, according to She et al. (2011) is that the given extracted constraints can be incomplete and, thus, not all relevant implications are contained in the implication graph.

2. After the feature hierarchy is specified, another auxiliary graph, a mutex graph, similar to the implication graph, is constructed. The mutex graph is an undirected graph with features as vertices and edges between two features u and v, if $u \implies \neg v$ and $v \implies \neg u$ hold with respect to the given constraints. That is, all adjacent features are mutually exclusive. Based on this mutex graph all maximal cliques (subsets of vertices that all are connected with each other) among the vertices

with the same parent are computed. All features within such a clique are mutually exclusive and share the same parent and represent mutex- or alternative-groups. She et al. (2011) introduce an additional constraint to extract xor-groups that require one of the groups' features to be selected if the parent is selected. This distinction is in line with the initial description of feature diagrams by Kang et al. (1990), but not all descriptions of feature diagrams follow this distinction between mutex- and xor-groups and just use the term alternative-group discussed in Sec. **??**.

3. CTCs for the feature diagram are extracted from the given configuration constraints. Since CTCs are constraints that could not be represented by the feature hierarchy (implication) or alternative-groups (exclusion) the derivation of CTCs follows this idea. The set of cross-tree implications is derived by removing all edges that are part of the feature hierarchy from the initially constructed implication graph. The set of cross-tree exclusions is derived in the same manner from the mutex graph by removing all edges among vertices of all mutex-groups. To make the feature model sound, the given configuration constraints, reduced to those clauses that are not already entailed by the diagram, can be added as an additional CTC formula to the feature diagram (She et al., 2011).

The approach by She et al. (2011) provides a practical algorithm to synthesize a feature diagram, yet has some aspects we might need to consider. First, the approach is not able to detect or-groups as defined in Sec. **??**. Second, the approach does introduce a root feature. Finally, the approach does not distinguish between mandatory and optional features. Implicitly, all features that do not have a parent feature are optional and all features that have a parent feature are by default mandatory. She et al. (2011) evaluated the algorithm with both complete and incomplete variability knowledge (feature names, descriptions and constraints). While the algorithm has shown to be practical, detecting features whose parent was the root-feature was difficult since, due to the transitive property of implication, it is implied by each feature of the feature model.

2.4   Assessing Performance

While the last three sections covered software evolution and variability modeling, we now step forward to the topic of software performance. This section will outline the term performance with respect to software systems as well as to possible measurements. We provide a brief look at the general performance testing setup and the required prerequisites, including suitable benchmarks. Finally, we provide the statistical background to summarize, interpret and compare performance measurements accurately.

2.4.1   What is Performance?

The performance of software systems is, like software quality, primarily a matter of perspective. While an end user might consider practical aspects to be more important, from a developer's perspective, performance relates to and is best described by non-functional properties (Molyneaux, 2014). While functional properties subsume what exactly a software system does, non-functional properties describe how (good or bad) a software system

is at providing the functionality offered. The notion of good and bad in this sense corresponds to non-functional requirements (NFR), that is, software with good performance behavior satisfies NFRs. The categories of NFRs that shape performance behavior are manifold. According to Molyneaux (2014), the categories or *key performance indicators* (KPIs) include

- availability

- response time,

- throughput,

- resource utilization, and in a broader scope also

- capacity.

Time-related KPIs are availability and response time, whereby availability describes the time or time ratio that the software is available to the end user, and response time subsumes the time it takes to finish a request or operation. Throughput as a category subsumes the program behavior with respect to program load, such as hits per second for a web application or amount of data processed per second. Resource utilization describes the extent to which a software system uses the physical resources (CPU time, memory, and disk or cache space) of the host machine. Finally, from a web-centered perspective, capacity describes measurements with respect to servers and networks, such as network utilization (traffic, transmission rate) and server utilization, such as resource limitations per application on a host server (Molyneaux, 2014).

Consequently, the assessment of performance requires a context or testing target that corresponds to the assessed system under test (SUT). For instance, for a simple command-line compression tool, suitable KPIs are response time and throughput, whereas performance for an online shop web application is better outlined by availability and capacity.

2.4.2 Performance Testing

The first step in performance testing, prior to defining relevant KPIs and metrics, is to specify a system operation or use case (Woodside et al., 2007) to assess performance for. A typical use case includes a well defined task or workload to process, expected behavior, outcome, and performance requirements as previously discussed. For the SUT, however, we require a version that does compile or, in case it is interpreted, is syntactically correct (Molyneaux, 2014). With regard to performance assessments as part of the development process, a code freeze should be obtained since measurement results are likely to become meaningless for later versions. In addition to that, the machine or setup used for performance measurement should ideally be as close to the production environment as possible, but at least be documented to compare different runs (Molyneaux, 2014).

Finally, one or more benchmarks need to be selected to simulate the program load for the respective use case. A benchmark, all in all, needs to be representative, i.e., should relate to the the use case or requirement one wants to validate. While benchmarks for file compression usually include multiple different types of media data (text, sound, pictures) like the Canterbury corpus[1], web applications can be exposed to dealing with a

---

[1]bener

number of simulated users at the same time ([Molyneaux, 2014](#)). Benchmarks have often been standardized within research and engineering communities to provide comparable performance measurements. A popular example is the Software Performance Evaluation Corporation (SPEC), a consortium providing a variety of benchmarks like the CPU2000[2] processor benchmark consisting of programs with both floating point and integer operations. Moreover, benchmarks in a sense of repeatable program load can be obtained from load generation software like ApacheBench[3] for the Apache web server or simply by measuring performance for test cases ([Heger et al., 2013](#); [Nguyen et al., 2014](#)).

Performance testing heavily relies on tool support, especially for repeating test cases and recording measurements. Since the tool solutions for performance testing vary from domain, scale and purpose, we will only outline the general tool architecture. [Molyneaux](#) [(2014)](#) describes four primary components for a performance testing setup: a scripting module, a test management module, a load injector, and an analysis module. A scripting module handles the generation or repetition of use cases which, for instance, can be recorded prior to the test for web applications. A test management module creates and executes a test session, whose program load is generated by one or more load injectors. A load injector can provide a benchmark by generating items to process, or can simulate a number of clients for a server-side application. An analysis module, finally, collects and summarizes data related to the performance testing target. More on summarizing and comparing recorded results in the next section.

2.5 Statistical Considerations

In this section we now discuss the appropriate statistical means to summarize, compare and interpret measurement results. For the remainder of this section we will refer to the following two use cases. First, to obtain robust results, the assessment of a single variant needs to be repeated multiple times. Consequently, to report a single result per metric, the measurements for a single variant need to be summarized. Second, the assessment of performance for a variant may comprise several use cases or system operations, for instance file compression and decompression for a compression software. Therefore, performance measures aggregated from different benchmarks need to be summarized accurately.

2.5.1 Measures of Central Tendency

For each test run of a software system or variant, we obtain a single-valued measurement per performance metric. Since we repeat each test run n times per variant, we obtain a data record X with n measurements $X = X_1, X_2, \ldots, X_{n-1}, X_n$. While the arithmetic mean is commonly considered the right means to summarize data records and report a representative average value, we need to be more cautious with how to summarize data records ([Fleming and Wallace, 1986](#); [Smith, 1988](#)). From a statistical perspective, the intention of summarizing a data record is to find a measure of central tendency what is representative for the data record. While the arithmetic mean is an appropriate method for many cases, there exist other means to summarize data records. Moreover, there exist a number of criteria for when to use which means to summarize a data record.

---

[2]The benchmark description can be found at `https://www.spec.org/cpu2000/`.

[3]Manual page of the Apache tool `ab`, `http://httpd.apache.org/docs/2.4/en/programs/ab.html`.

The first question when summarizing a data record is to ask what the data actually describe and what we to express with our summary. For a data record X, we can define a relationship we would like to keep while replacing each single X_i with an average value $\bar{x}$. Based on this relationship, we can derive the appropriate method to summarize our record. For instance, our data record X describes the time elapsed for a test case and we want the keep the following relation, saying that the sum of all measurements $\sum_{i=1}^{n}$ is defines the total time elapsed T as

$$T = \sum_{i=1}^{n} X_i = \sum_{i=1}^{n} \bar{x}.$$

Based on the term above, we can derive the definition of the method appropriate to summarize our data record with respect to the conserved relation what is the *arithmetic mean*, defined as

$$\bar{x} = \frac{1}{n} \sum_{i=1}^{n} X_i. \tag{2.1}$$

Consider another example, similar to the one above, where the test case is a load test with a predefined number of users and the measurements $X = X_1, X_2, \ldots, X_{n-1}, X_n$ are measured as hits per second. Again, we have different measurements we want to summarize with respect to a relation to conserve. Each user drops the same number of requests T, whereby the hit rate X_i and the respective elapsed time $t_i = \frac{T}{X_i}$ vary. We now conserve the relationship that the total number of requests for the data record is the sum of all hit rates X_i times the time elapsed t_i. Therefore, we want to substitute each hit rate X_i with an average value $\hat{x}$ so that the aforementioned rule holds

$$n \cdot T - \sum_{n}^{i=1} X_i t_i = 0 \Leftrightarrow n \cdot T - \hat{v} \sum_{n}^{i=1} t_i = 0 \Leftrightarrow n = \hat{x} \sum_{n}^{i=1} \frac{1}{X_i}$$

Similar to Eq. 2.1 we can derive the definition for the summarization method to use from the equation above what is the *harmonic mean*, defined as

$$\hat{x} = n \cdot \left(\sum_{i=1}^{n} \frac{1}{X_i} \right)^{-1} = \frac{n}{\frac{1}{X_1} + \frac{1}{X_2} + \ldots + \frac{1}{X_{n-1}} + \frac{1}{X_n}}. \tag{2.2}$$

We see that the summarization methods presented in Eq. 2.1 and 2.2 are useful for different types of data records and they should be used accordingly. The arithmetic mean is suitable for records, where the total sum of single measurements has a meaning, whereas the harmonic mean is suitable for measured rates or frequencies (Smith, 1988).

While the aforementioned measures of central tendency should be appropriate for most cases, they are not always the best choice though. The arithmetic and harmonic mean are heavily influenced by extreme observations (Shanmugam and Chattamvelli, 2015). For instance, for the data record $X = 1, 2, 3, 30$ three of four values are smaller than the arithmetic mean $\bar{x} = 9$. One option is to explicitly exclude outlier values from the data record. The so-called *trimmed mean* is obtained by truncating a upper and/or lower percentage t of the data record and, consequently, computing the (arithmetic or harmonic) mean for the remaining data record (Shanmugam and Chattamvelli, 2015).

While this method is suitable to omit the effects of outliers, one still needs to specify which upper and/or lower percentage t needs to be truncated. Moreover, the use a (trimmed) mean requires the data records' frequencies to be distributed symmetrically around its mean. A probability distribution is skewed (and therefore asymmetric) if, graphically speaking, its histogram is not symmetric around its measure of central tendency. A simple method to measure the skewness of a probability distribution is *Bowley's measure* (Shanmugam and Chattamvelli, 2015), defined as

$$B_S = \frac{(Q_3 - M) - (M - Q_1)}{Q_3 - Q_1} = \frac{(Q_3 + Q_1 - 2M)}{Q_3 - Q_1}, \tag{2.3}$$

where Q_1 and Q_3 denote the first and third quartile, and M denotes the median of the probability distribution. The quartiles Q_i with $i \in \{1, 2, 3\}$ are defined as the values of a data record X, so that $\frac{i}{4}$ of the values of X are smaller than Q_i. The *median* is defined as Q_2, i.e., a value $M \in X$ so that half of the values in X are smaller than M.

The median itself is a more robust measure of central tendency than the aforementioned ones since it is less influenced by outliers and can be used for both skewed and symmetric data records (Shanmugam and Chattamvelli, 2015). For a given ascendingly ordered data record $X = X_1, X_2, \ldots, X_{n-1}, X_n$, we can compute the median as follows:

$$\text{Median}(X) = \begin{cases} X_{\frac{n+1}{2}} & \text{if } n \text{ is odd} \\ \frac{1}{2}\left(X_{\frac{n}{2}} + X_{\frac{n}{2}+1}\right) & \text{if } n \text{ is even} \end{cases} \quad \text{with } X_i \leq X_j \text{ and } i < j \leq n \tag{2.4}$$

2.5.2 Measures of Dispersion

Last, we take a look at different measures of spread or dispersion. Most commonly used are the *variance* σ^2 and the *standard deviation* $\sigma = \sqrt{\sigma^2}$ defined along the mean μ of a probability distribution as

$$\sigma^2 = \frac{\sum_{i=1}^{n}(X_i - \mu)^2}{n} \tag{2.5}$$

Similar to the (arithmetic or harmonic) mean, the variance and the standard deviation are heavily influenced by extreme observations (Shanmugam and Chattamvelli, 2015) are not the best choice in all cases. Instead, two more robust measures are the *median absolute deviation* (MAD) and the *inter-quartile range* (IQR). The MAD is defined as the median of the absolute deviations from the probability distributions' median (Molyneaux, 2014), or, defined as follows:

$$\text{MAD}(X) = \text{Median}\Big(|X - \text{Median}(X)|\Big) \tag{2.6}$$

The IQR, however, defines the range between the first and the third quartile Q_1 and Q_3 (Shanmugam and Chattamvelli, 2015). This range is larger for a widespread data record and smaller for a data range with a narrow spread, but is not influenced by extreme observations as those outliers do not lie within the range $[Q_1, Q_3]$. The IQR is defined as follows

$$\mathrm{IQR}(X) = Q_3 - Q_1. \tag{2.7}$$

2.5.3   When to use which measure?

In this section we discussed the arithmetic and harmonic mean as well as the median as a measure of central tendency, the symmetric property that is required to use the arithmetic or harmonic mean, and different measures of spread for a given data record. At the beginning we asked for means to summarize data records a) for an experiment repeated multiple times, and b) obtained from different benchmarks. While for case b) the answer is to use the arithmetic or harmonic mean (depending on the quality of the measurements), for a) the answer is a little more elaborate.

The arithmetic (or harmonic) mean can be used whenever the quality of the measurement is appropriate and the data record is symmetric. According to Shanmugam and Chattamvelli (2015), the arithmetic mean is preferable "when the numbers combine additively to produce a resultant value", such as time periods or memory sizes, whereas the harmonic mean is preferable "when reciprocals of several nonzero numbers combine additively to produce a resultant value", such as rates or frequencies (Smith, 1988). The median is a less precise estimator than the both means, yet more robust with regard to extreme observations. In addition, the MAD or IQR provide a more robust means of spread than the standard deviation.

2.6   Performance Prediction Models

In the last section we referred to performance, or in detail, the KPIs, as possible testing targets we validate against non-functional requirements. However, in a broader sense, software performance has become an aspect of software engineering referred to as *software performance engineering* (SPE) (Woodside et al., 2007). A lot of effort has been spent to study and describe performance behavior as well as to improve performance quality. Besides the analysis of concerns or requirements with respect to performance, SPE comprises performance testing as well as performance prediction (Woodside et al., 2007). Performance prediction aims at modeling and estimating performance behavior for different use cases or configurations. The first approaches to performance prediction models model the underlying software system component or operation from which performance estimations are then deduced. These so-called *model-based prediction models* enable a performance estimation early in the development process since no actual performance measurement is required (Woodside et al., 2007). In opposition to model-based approaches, measurement-based approaches have emerged. These *measurement-based prediction models* are based on a sample of performance measures which are used to learn a software systems' performance behavior (Woodside et al., 2007). More recently, measurement-based prediction models that emphasize variability have been proposed, of which we will discuss three approaches in the remainder of this section.

Learning Performance   In essence, learning and predicting performance behavior for a configurable system means nothing different than finding an approximation h for a function $f : C \to \mathbb{R}$ where C is the set of configurations and $f(c) \in \mathbb{R}$ with $c \in C$ is the corre-

sponding performance measurement. The accuracy of the approximation $\hat{f}$ describes how the estimated performance $\hat{f}(c)$ deviates from the actually measured performance $f(c)$. Different approaches to construct such an approximation, referred to as a *performance prediction model*, emerged recently. All approaches utilize two samples of configurations and corresponding performance measurements to build and validate the model. A training sample is used to learn the approximation from, whereas a testing sample is used to assess the prediction error rate of the previously learned approximation.

A straightforward methodology to learn performance behavior was proposed by Guo et al. (2013). The authors utilize classification-and-regression-trees (CART), akin to decision trees, to derive a top-down classification hierarchy for a given sample. The approach supports progressive (and random) sampling, i.e., the performance model is constructed several times while the size of the training sample is successively increased. The model construction commences with a recursive segmentation to find an an accurate CART for a local subset of the training sample. Subsequently, the local trees are merged to constitute a CART for the whole training sample. It is worth mentioning that the estimation using CART is not limited to binary configuration options, but supports numeric features as well. Moreover, the approach by Guo et al. (2013) does not produce any further computation overhead besides the measurement effort and the construction of a CART.

A different approach to modeling performance behavior was proposed by Zhang et al. (2015). The authors propose an approach to construct performance prediction models based on a Fourier description of the performance-describing function f. The principal idea is to approximate a Fourier series approximating the function performance-describing f, and learning all coefficients of the series terms. The number of terms is exponential in the number of configuration options. The main characteristic of this approach is that, prior to learning a performance prediction model, a desired level of prediction accuracy can be specified. That is, the approach automatically chooses the sample size required to learn the prediction model accordingly.

A third approach to model and predict performance behavior, *SPLConqueror*, was proposed by Siegmund et al. (2012). The authors describe performance behavior for a configuration as the accumulated influence of features the configuration is composed from, and respective feature interactions. To estimate the influence of single features and feature interactions, the authors propose a number of sampling strategies. Single features are assessed by comparing the performance of two different configurations per feature: For a feature f, two valid configurations are compared, whereby both configurations have the minimal number of features selected that are not excluded by the selection of f. While for one configuration f is selected, it is deselected for the other one. The difference between the performance measures for both configurations is the estimate $\Delta_{min(f)}$ for the performance influence of feature f. Feature interactions that are performance-relevant are detected automatically in a similar manner. The main idea is, that a feature f is more likely to interact with other features, the more features are selected along with f. Therefore, if f interacts, the influence of feature f, $\Delta_{min(f)}$ differs from the influence of f, when estimated using configurations where the maximal number of features selectable together with f are used. Based on the set of interacting features, three heuristics are employed to detect feature interactions. Simple feature interactions are detected using pair-wise sampling (Siegmund et al., 2012; Apel et al., 2013), whereas higher-order feature interactions are often composed from simpler ones, or centered around a small subset of hot-spot features (Siegmund et al., 2012). Finally, for an arbitrary configuration the

performance can be predicted by the sum of influence estimations per feature and feature interaction.

The idea of predicting performance by estimating the influence of feature and feature interactions was further developed by Siegmund et al. (2015) by proposing performance influence models for both binary and numeric configuration options. A performance influence, similar to *SPLConqueror*, predicts performance by computing the sum of previously learned influence estimates. The novelty of this approach lies in the way the terms describing the influence of feature and feature interactions are learned and not derived from a sampling strategy. Instead of a single constant performance influence estimate, each term can contain of a number of functions, such as linear, quadratic or logarithm functions, or compositions thereof. This does not only allow to incorporate domain knowledge by preselecting functions. Moreover, by introducing combinations of non-linear functions and learning the coefficients using linear regressions enables learning a non-linear function. The algorithm commences with the selection of terms and successively extends and reduces the term selection to increase the prediction accuracy. The outcome, similar to Siegmund et al. (2012), is a linear function whose terms represent the estimated influence of features and feature interactions on performance.

All aforementioned approaches rely on performance measurements and, to some extent, also on a sampling strategy. While the approaches by Siegmund et al. (2012, 2015) essentially describe sampling strategies, the approach proposed by Zhang et al. (2015) is able to work with random samples, and the approach of Guo et al. (2013) does extend its training sample automatically. As performance measurements entails an expensive and time-consuming process, Sarkar et al. (2015) have investigated different sampling strategies in this domain. The authors compare progressive sampling, where sample sizes are increased successively until the learned prediction model performs accurately enough, and projective sampling, where the optimal sample size is estimated based on the expected learning curve, with respect to cost and prediction accuracy. They advocate the use of projective sampling over progressive sampling. In addition, the authors propose and evaluate an own sampling heuristic to select an initial sample for progressive sampling. This sampling heuristic is based on feature frequencies and outperforms t-wise sampling, such as pair-wise sampling (Sarkar et al., 2015).

3. Performance Measurement Setup

3.1 Performance Measurement Infrastructure

3.1.1 Repository Retrieval

3.1.2 Configuration Generation

Prior to any test session, we require a set of valid configurations for the configurable software system to benchmark. The configuraion semantics is commonly expressed in feature diagrams (cf. Sect. 2), but can be translated to and from propositional formulas or context-free grammars (CFGs) (although with additional CTCs) as well (Batory, 2005). The next section presents a methodology to translate a given feature diagram to an equivalent CFG, derive valid words from the CFG, and subsequently transform derived words to configuration artifacts.

Def. 3.1.1 (Context-free grammar). *A context-free grammar is a tuple $G = (N, T, S, P)$ with a set of non-terminal symbols N, terminal symbols T, a start word $S \in (N \cup T)^*$, and set of productions $P \subseteq N \times (N \cup T)^*$.*

Def. 3.1.2 (Configuration grammar). *A configuration grammar (CG) is a CFG, whereby all features represent a non-terminal symbols. For each binary feature $b \in N$, we introduce two terminal symbols $(b, 0)$ and $(b, 1)$. For each numeric feature $n \in N$, we introduce as many terminal symbols of the form (n, v) as there are valid numeric values for n, where, $v \in dom(n)$. Based on the feature hierarchy represented by the underlying feature diagram, we introduce a production in the following manner. The set of productions is*

$$P_{CG} = \{(p, \{c\}) \mid p, c \in N \land p \implies c\}$$

Data: this text
Result: how to write algorithm with LaTeX2e
initialization;
while *not at end of this document* **do**
    read current;
    if *understand* **then**
        go to next section;
        current section becomes this one;
    else
        go back to the beginning of current section;
    end
end

Algorithm 1: How to write algorithms

3.1.3   Sampling strategies

3.1.4   Benchmarks

3.1.5   Measurements with `GNU time`

3.1.6   Measurement Aggregation

3.2   Experiment Optimizations

3.3   Preprocessing Strategies

3.3.1   Variabilty Model Extraction

3.3.2   Build Mechnism Extraction

3.3.3   Release Commit Extraction

3.4   Case Study Corpus

4. Case Study Evaluation

5. Conclusion

Bibliography

Alves, V., C. Schwanninger, L. Barbosa, A. Rashid, P. Sawyer, P. Rayson, C. Pohl, and A. Rummler
2008. An exploratory study of information retrieval techniques in domain analysis. In *Software Product Line Conference, 2008. SPLC'08. 12th International*, Pp. 67–76. IEEE.
Cited on pages 4 and 9.

Andersen, N., K. Czarnecki, S. She, and A. Wąsowski
2012. Efficient synthesis of feature models. In *Proceedings of the 16th International Software Product Line Conference-Volume 1*, Pp. 106–115. ACM.
Cited on pages 4 and 11.

Apel, S., D. Batory, C. Kästner, and G. Saake
2013. *Feature-Oriented Software Product Lines*. Berlin, Heidelberg: Springer Berlin Heidelberg. DOI: 10.1007/978-3-642-37521-7.
Cited on pages 1, 2, 3, 6, and 18.

Bakar, N. H., Z. M. Kasirun, and N. Salleh
2015. Feature extraction approaches from natural language requirements for reuse in software product lines: A systematic literature review. *Journal of Systems and Software*, 106:132–149.
Cited on pages 4 and 9.

Batory, D.
2005. Feature models, grammars, and propositional formulas. In *SPLC*, volume 3714, Pp. 7–20. Springer.
Cited on pages 6 and 20.

Breivold, H. P., I. Crnkovic, and M. Larsson
2012. A systematic review of software architecture evolution research. *Information and Software Technology*, 54(1):16–40.
Cited on page 8.

Czarnecki, K. and U. W. Eisenecker
2000. *Generative Programming: Methods, Tools, and Applications*. ACM Press/Addison-Wesley Publishing Co.
Cited on page 6.

Dietrich, C., R. Tartler, W. Schröder-Preikschat, and D. Lohmann
2012. A robust approach for variability extraction from the Linux build system. In *Pro-*

ceedings of the 16th International Software Product Line Conference-Volume 1, Pp. 21–30. ACM.
Cited on page 1.

Fleming, P. J. and J. J. Wallace
1986. How not to lie with statistics: the correct way to summarize benchmark results. *Communications of the ACM*, 29(3):218–221.
Cited on page 14.

Guo, J., K. Czarnecki, S. Apely, N. Siegmundy, and A. Wasowski
2013. Variability-aware performance prediction: A statistical learning approach. In *Proceedings of the 28th IEEE/ACM International Conference on Automated Software Engineering*, Pp. 301–311. IEEE Press.
Cited on pages 4, 18, and 19.

Guo, Y., C. Seaman, R. Gomes, A. Cavalcanti, G. Tonin, F. Q. B. Da Silva, A. L. M. Santos, and C. Siebra
2011. Tracking technical debt — An exploratory case study. Pp. 528–531. IEEE.
Cited on page 3.

Heger, C., J. Happe, and R. Farahbod
2013. Automated root cause isolation of performance regressions during software development. In *Proceedings of the 4th ACM/SPEC International Conference on Performance Engineering*, Pp. 27–38. ACM.
Cited on pages 3 and 14.

Hunsen, C., B. Zhang, J. Siegmund, C. Kästner, O. Le\s senich, M. Becker, and S. Apel
2016. Preprocessor-based variability in open-source and industrial software systems: An empirical study. *Empirical Software Engineering*, 21(2):449–482.
Cited on page 1.

Kang, K. C., S. G. Cohen, J. A. Hess, W. E. Novak, and A. S. Peterson
1990. Feature-oriented domain analysis (FODA) feasibility study. Technical report, Carnegie-Mellon Univ Pittsburgh Pa Software Engineering Inst.
Cited on pages 6, 11, and 12.

Molyneaux, I.
2014. *The Art of Application Performance Testing: Help for Programmers and Quality Assurance*, 2 edition. O'Reilly Media, Inc.
Cited on pages 2, 3, 12, 13, 14, and 16.

Nadi, S., T. Berger, C. Kästner, and K. Czarnecki
2014. Mining configuration constraints: Static analyses and empirical results. In *Proceedings of the 36th International Conference on Software Engineering*, Pp. 140–151. ACM.
Cited on page 10.

Nadi, S., T. Berger, C. Kästner, and K. Czarnecki
2015. Where do configuration constraints stem from? an extraction approach and an

empirical study. *IEEE Transactions on Software Engineering*, 41(8):820–841.
Cited on pages 4, 9, and 10.

Nguyen, T. H., M. Nagappan, A. E. Hassan, M. Nasser, and P. Flora
2014. An industrial case study of automatically identifying performance regression-causes. In *Proceedings of the 11th Working Conference on Mining Software Repositories*, Pp. 232–241. ACM.
Cited on pages 3 and 14.

Parnas, D. L.
1994. Software aging. In *Proceedings of the 16th international conference on Software engineering*, Pp. 279–287. IEEE Computer Society Press.
Cited on pages 7 and 8.

Passos, L., K. Czarnecki, and A. Wąsowski
2012. Towards a catalog of variability evolution patterns: the Linux kernel case. In *Proceedings of the 4th International Workshop on Feature-Oriented Software Development*, Pp. 62–69. ACM.
Cited on pages 2 and 9.

Passos, L., J. Padilla, T. Berger, S. Apel, K. Czarnecki, and M. T. Valente
2015. Feature scattering in the large: a longitudinal study of Linux kernel device drivers. Pp. 81–92. ACM Press.
Cited on pages 2 and 8.

Peng, X., Y. Yu, and W. Zhao
2011. Analyzing evolution of variability in a software product line: From contexts and requirements to features. *Information and Software Technology*, 53(7):707–721.
Cited on pages 2, 3, 4, 8, and 9.

Perry, D. E. and A. L. Wolf
1991. Software architecture. *Submitted for publication.*
Cited on pages 2 and 8.

Rabkin, A. and R. Katz
2011. Static extraction of program configuration options. In *Proceedings of the 33rd International Conference on Software Engineering*, Pp. 131–140. ACM.
Cited on pages 4, 9, and 10.

Sarkar, A., J. Guo, N. Siegmund, S. Apel, and K. Czarnecki
2015. Cost-efficient sampling for performance prediction of configurable systems (t). In *Automated Software Engineering (ASE), 2015 30th IEEE/ACM International Conference on*, Pp. 342–352. IEEE.
Cited on page 19.

Schaefer, I., L. Bettini, V. Bono, F. Damiani, and N. Tanzarella
2010. Delta-oriented programming of software product lines. *Software Product Lines: Going Beyond*, Pp. 77–91.
Cited on page 1.

Seidl, C., F. Heidenreich, and U. A\s smann
2012. Co-evolution of models and feature mapping in software product lines. In *Proceedings of the 16th International Software Product Line Conference-Volume 1*, Pp. 76–85. ACM.
Cited on pages 2 and 9.

Shanmugam, R. and R. Chattamvelli
2015. *Statistics for scientists and engineers.* Hoboken, New Jersey: Wiley.
Cited on pages 15, 16, and 17.

She, S., R. Lotufo, T. Berger, A. Wasowski, and K. Czarnecki
2011. Reverse engineering feature models. In *Software Engineering (ICSE), 2011 33rd International Conference on*, Pp. 461–470. IEEE.
Cited on pages 4, 11, and 12.

Siegmund, N., A. Grebhahn, S. Apel, and C. Kästner
2015. Performance-influence models for highly configurable systems. Pp. 284–294. ACM Press.
Cited on pages 1, 4, and 19.

Siegmund, N., S. S. Kolesnikov, C. Kästner, S. Apel, D. Batory, M. Rosenmüller, and G. Saake
2012. Predicting performance via automated feature-interaction detection. In *Software Engineering (ICSE), 2012 34th International Conference on*, Pp. 167–177. IEEE.
Cited on pages 4, 18, and 19.

Smith, J. E.
1988. Characterizing computer performance with a single number. *Communications of the ACM*, 31(10):1202–1206.
Cited on pages 14, 15, and 17.

Thüm, T., S. Apel, C. Kästner, I. Schaefer, and G. Saake
2014. A Classification and Survey of Analysis Strategies for Software Product Lines. *ACM Computing Surveys*, 47(1):1–45.
Cited on page 1.

White, J., B. Dougherty, and D. C. Schmidt
2009. Selecting highly optimal architectural feature sets with filtered cartesian flattening. *Journal of Systems and Software*, 82(8):1268–1284.
Cited on page 1.

Woodside, M., G. Franks, and D. C. Petriu
2007. The future of software performance engineering. In *Future of Software Engineering, 2007. FOSE'07*, Pp. 171–187. IEEE.
Cited on pages 3, 13, and 17.

Xu, T., L. Jin, X. Fan, Y. Zhou, S. Pasupathy, and R. Talwadker
2015. Hey, you have given me too many knobs!: understanding and dealing with over-designed configuration in system software. Pp. 307–319. ACM Press.
Cited on page 9.

Zhang, B., M. Becker, T. Patzke, K. Sierszecki, and J. E. Savolainen
2013. Variability evolution and erosion in industrial product lines: a case study. In *Proceedings of the 17th International Software Product Line Conference*, Pp. 168–177. ACM.
Cited on pages 2 and 8.

Zhang, Y., J. Guo, E. Blais, and K. Czarnecki
2015. Performance prediction of configurable software systems by fourier learning (T). In *Automated Software Engineering (ASE), 2015 30th IEEE/ACM International Conference on*, Pp. 365–373. IEEE.
Cited on pages 4, 18, and 19.

Zhou, S., J. Al-Kofahi, T. N. Nguyen, C. Kästner, and S. Nadi
2015. Extracting configuration knowledge from build files with symbolic analysis. In *Proceedings of the Third International Workshop on Release Engineering*, Pp. 20–23. IEEE Press.
Cited on pages 4 and 10.